

# Introduction to Service Mesh on Kubernetes

**Tran Trung Nguyen**

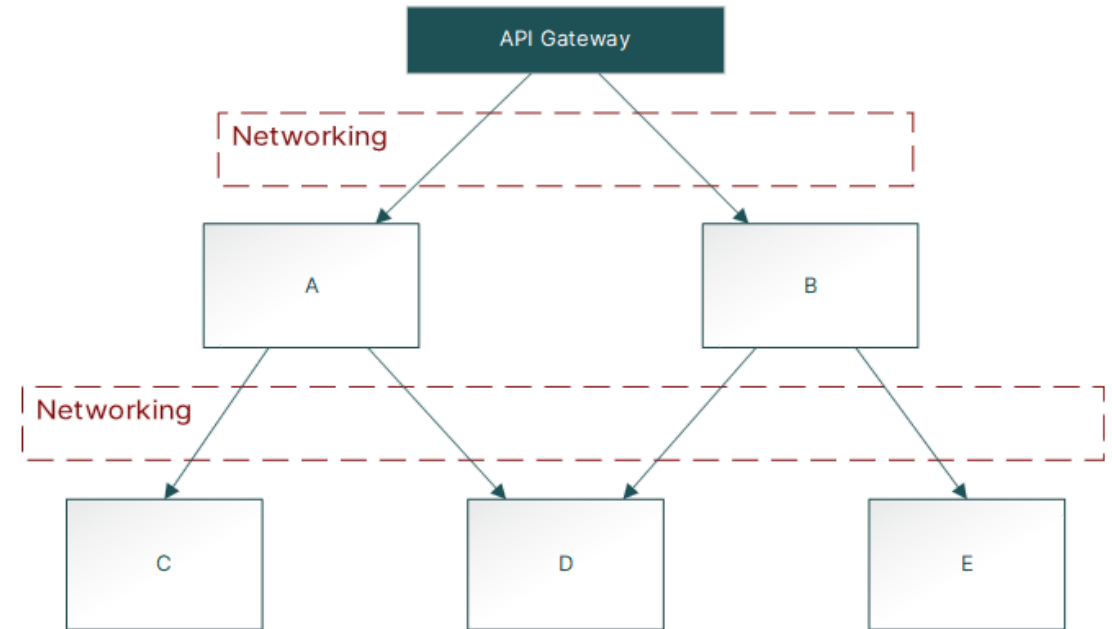
*Software Engineer*

VNPAY Cloud

# Microservices

are essentially distributed systems,  
often needing:

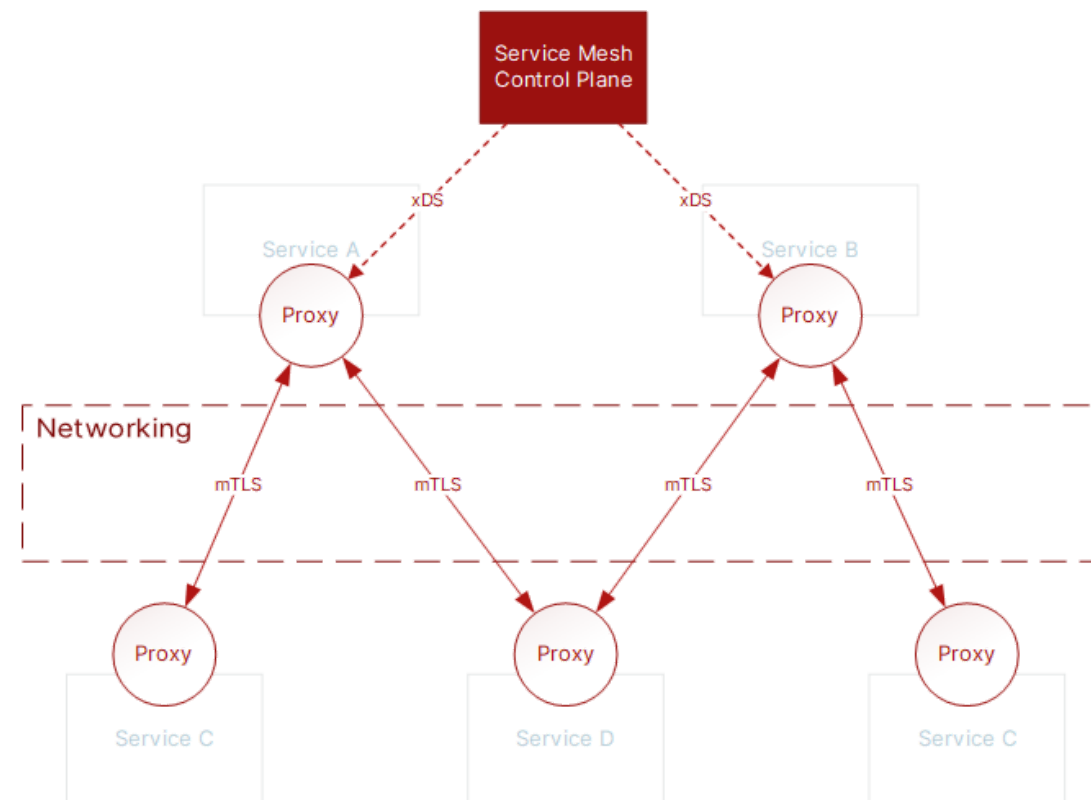
- Routing and Discovery
- Failure Handling (*Retries, Circuit Breaking, Timeout*)
- Encryption
- Observability (*Logging, Tracing, Metrics*)



# Service Mesh

Networking between microservices is considered a dedicated infrastructure layer, called *service mesh*

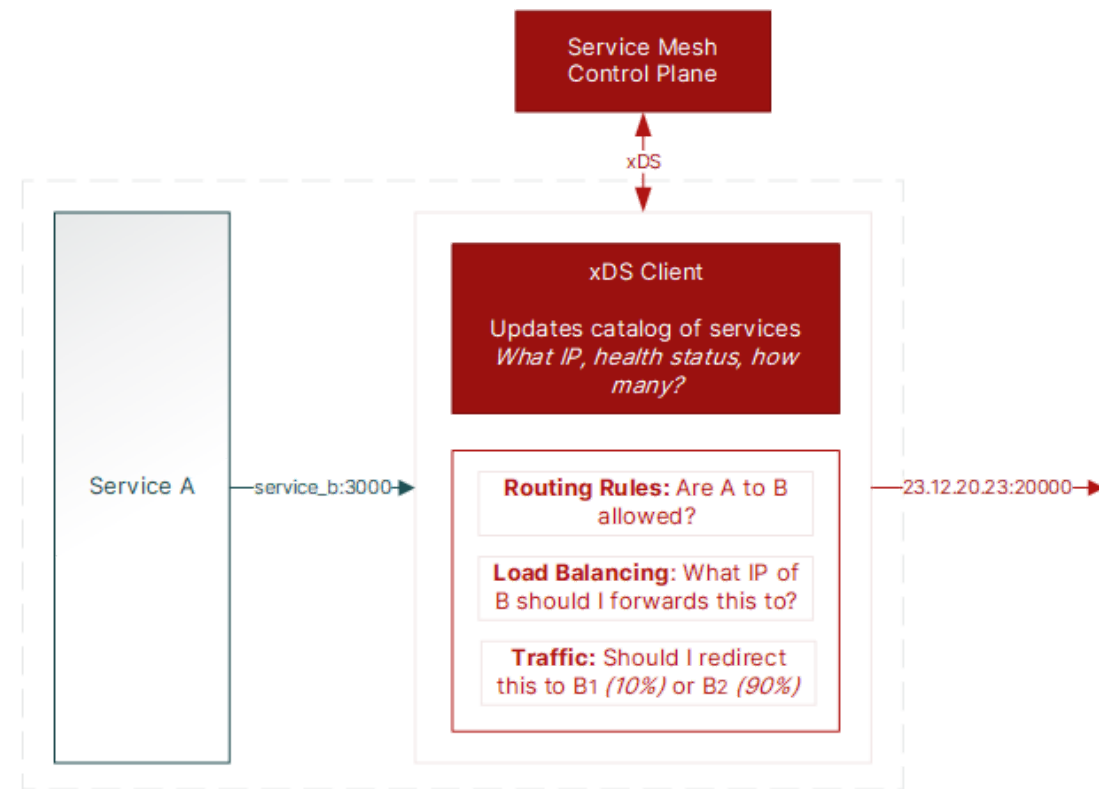
- *Mesh*: Discover and connects services
- Uses sidecar proxies to intercept and forwards communication



## Routing and Discovery

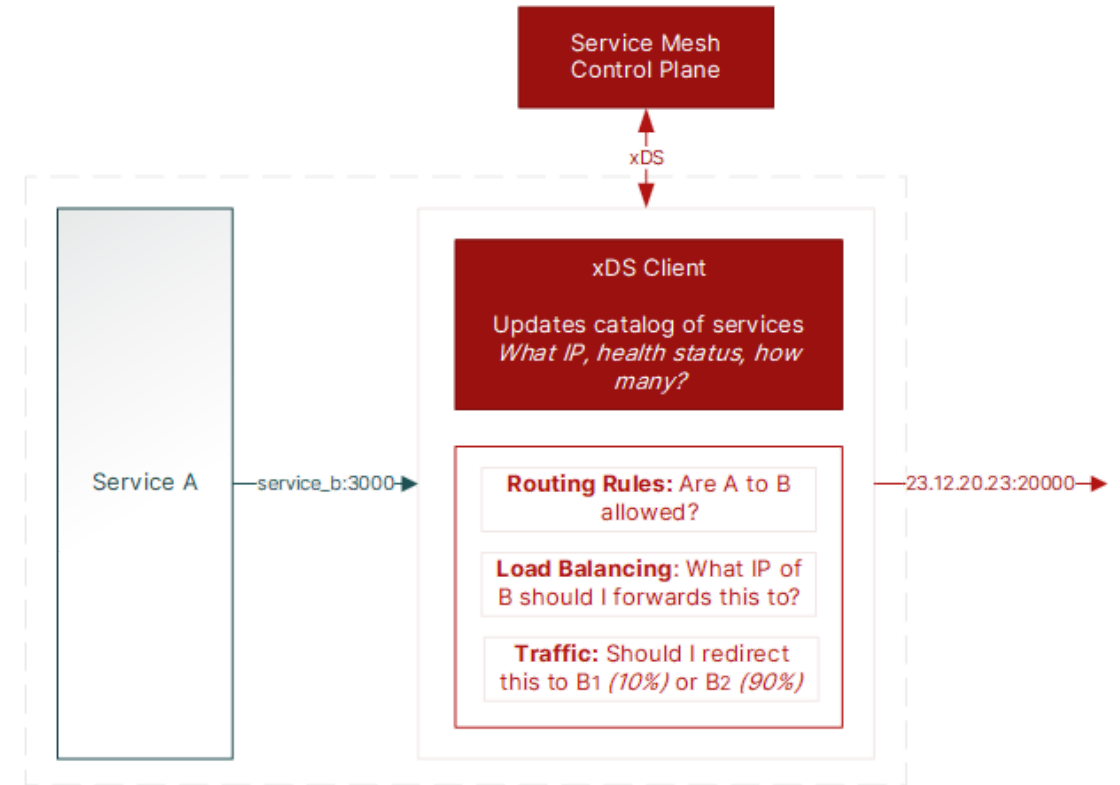
The application need only the name of other services. The *service mesh* do:

- *Discovery*: Updates from and to the control plane the status, IP address, etc. of the services
- *Routing Rules*: Service A is allowed to talk to B through gRPC, HTTP GETs



## Routing and Discovery

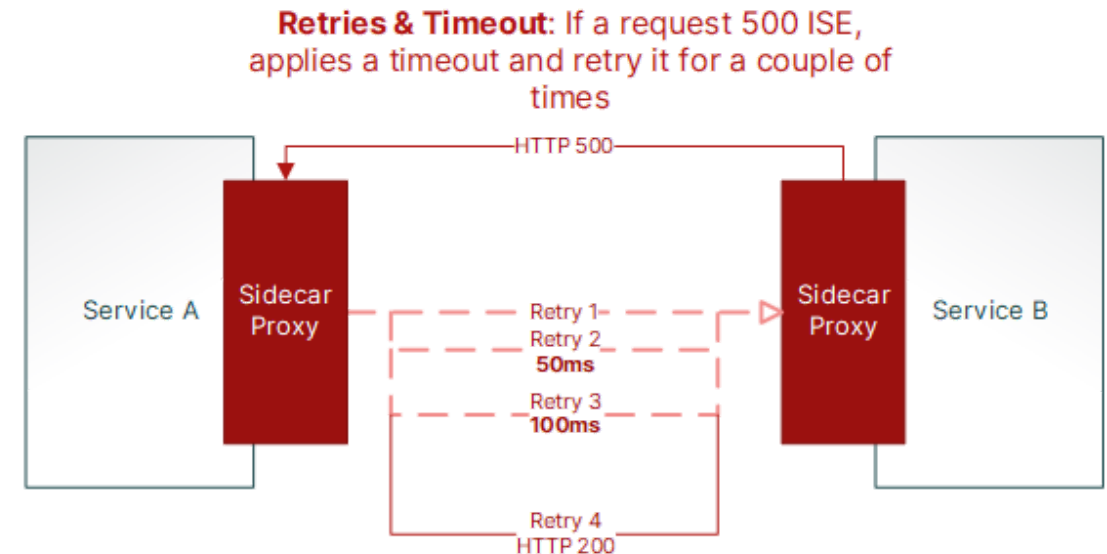

- *Load Balancing*: Kubernetes Service API randomly selects a Pod. Proxies often provides **round-robin, random, weighted, least request,...**
- *Traffic Splitting*: Great for A/B Testing. Can use probability or HTTP headers, paths, etc.



## Failure Handling

Often these features are built into application code. The *service mesh* can **automatically** do:

- *Circuit Breaking*: If a request to a service fails, stop the connection and gradually open it again until that service is back
- *Retries & Timeout*

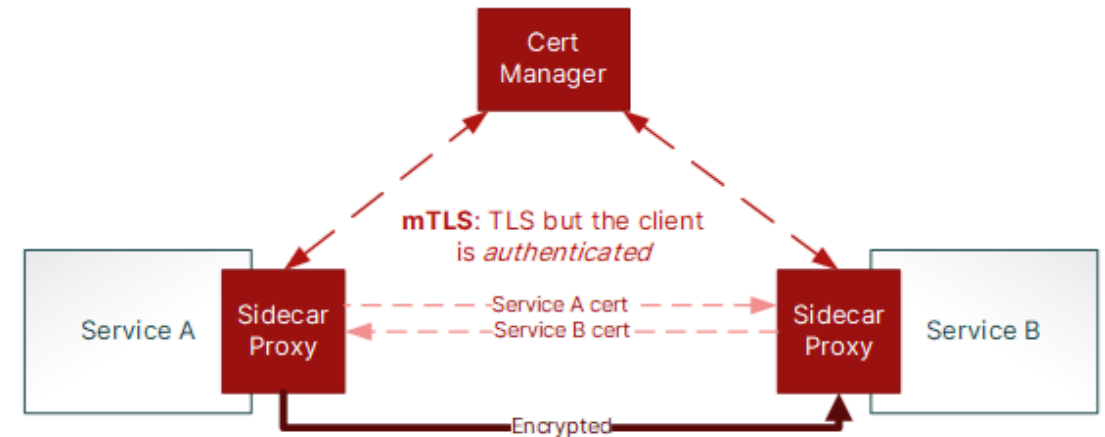


## Encryption

Either implement encryption from application-level code with **TLS** or use service mesh

- Service meshes often use **mTLS**, a form of TLS where the client also sends a client certificate
- Server checks the certificate identity and decide to accept handshake or not

Why not just use regular TLS?



# Control

**mTLS:** Sometimes it is useful to know who the client is

Knowing the client can help service mesh:

- *Control:* Is client A allowed to communicate with server B?  
**Example:** Prometheus should only use the `/metrics` API
- *Tracing and Metrics:* We know that client A is communicating with server B

The screenshot shows the 'New Intention' form in the Istio UI. The browser address bar shows 'localhost:8500/ui/dc1/intentions/create'. The navigation bar includes 'dc1', 'Services', 'Nodes', 'Key/Value', 'ACL', and 'Intentions'. The form has two dropdown menus: 'Source Service' with 'app-ui' selected and 'Destination Service' with 'news-service' selected. Below these are three radio button options for 'Should this source connect to the destination?': 'Allow' (selected), 'Deny', and 'Application Aware'. Each option has a brief description. At the bottom, there is an optional 'Description' text area and 'Save' and 'Cancel' buttons.

localhost:8500/ui/dc1/intentions/create

dc1 Services Nodes Key/Value ACL Intentions

< All Intentions

### New Intention

**Source**  
Source Service  
app-ui  
Search for an existing service, or enter any Service name.

**Destination**  
Destination Service  
news-service  
Search for an existing service, or enter any Service name.

Should this source connect to the destination?

☒ **Allow**  
The source service will be allowed to connect to the destination.

☐ **Deny**  
The source service will not be allowed to connect to the destination.

☐ **Application Aware**  
The source service may or may not connect to the destination service via unique permissions based on Layer 7 criteria: path, header, or method.

**Description (Optional)**  
Description (Optional)


Save Cancel



## Observability

Collects the 4 **Golden Signals** with ease:

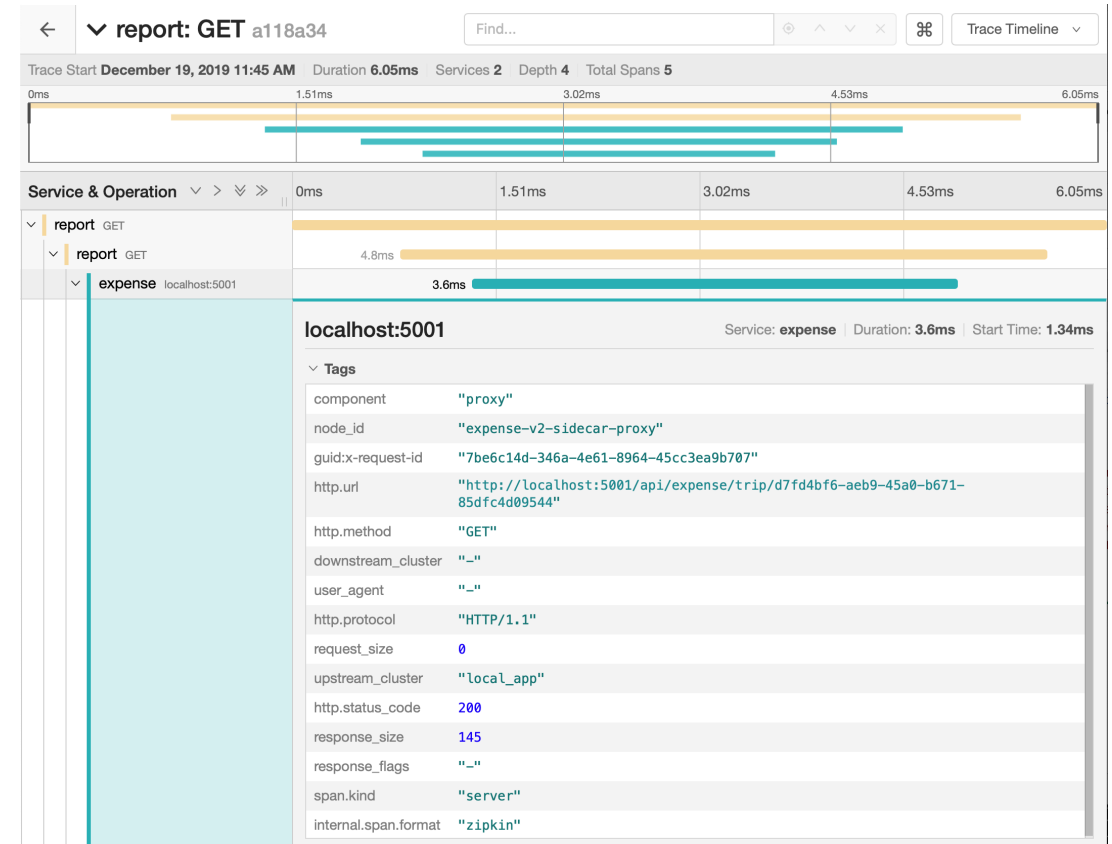
1. *Latency*: The time it takes to service a request
2. *Traffic*: How much demand is being placed on the system
3. *Errors*: How often the requests fail
4. *Saturation*: How much more traffic can the services handle?



# Observability

Most service meshes provide means to:

- *Collect metrics*: Envoy proxies can record latency, request rates and error rates. Use other ways to collect resource usage.
- *Export metrics*: Most exports to Prometheus and provide a built-in dashboard to displays the metrics. Some supports exporting metrics to distributed tracing platforms



## Options

There are plenty of Service Mesh implementations to choose from:

- **CNCF Projects:** Istio, Linkerd, Kuma, Open Service Mesh,...
- **HashiCorp:** Consul
- **Others:** Traefik Mesh
- **AWS:** App Mesh

For more, visit: [servicemesh.es](https://servicemesh.es)




# Demo

Service Mesh with Consul on Kubernetes

## The problem with sidecar proxies

Sidecar proxies costs:

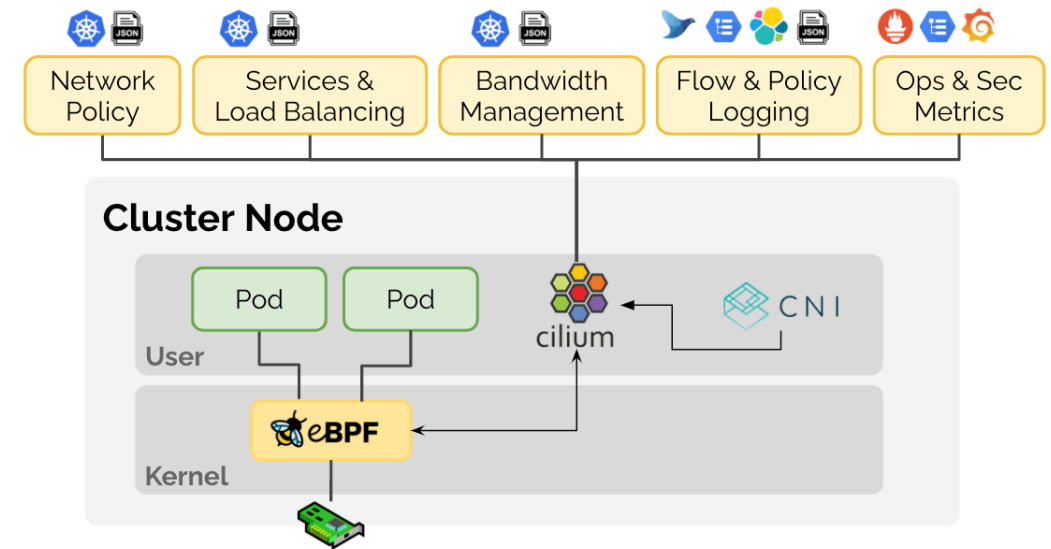
- *Resources*: Each pod requires an Envoy proxy, consuming ~200 CPU units with ~40-50MB of RAM
- *Performance*: Extra network hop means extra delay; slower pod start up time
- *Effort*: Maintaining the connections to the control plane requires scaling; debugging becomes harder; another dependency



# Cilium

Started out as a CNI (*Container Network Interface*) for Kubernetes

- *Connects services*: IPv4, IPv6, etc.
- *Enhance security*: Kubernetes Network Policy, Transparent Encryption
- *Load balancing*: Replacement for *kube-proxy*
- *Observability*: Network flow logs, metrics, L3/4 & L7 (*not really*) monitoring



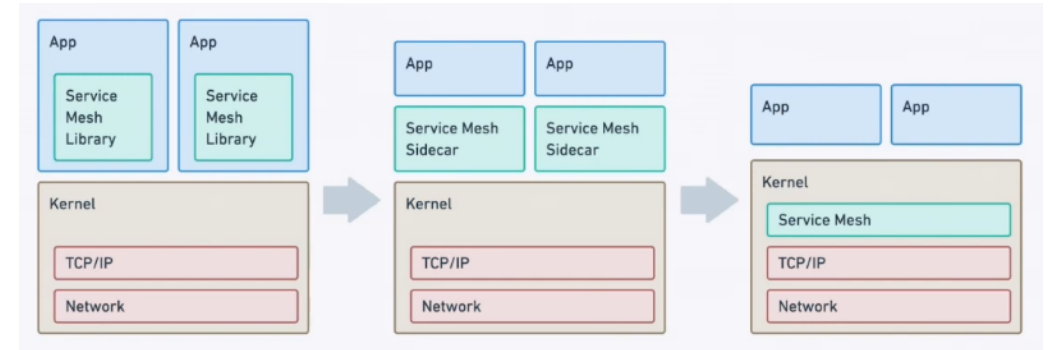
## eBPF

Makes kernel **extensible**.

It allows access to kernel-level APIs, bringing better **networking visibility**

- Better networking visibility allows more **control**
- Sits between Pods and forwards network packets to Pod network namespaces based on policies, encrypting it along the way


**YouTube:** [Effortless Mutual Authentication with Cilium](#)



## Do you need service mesh?

It depends, but if you need...

- **Need pure L3/L4 networking administration:** Use Kubernetes CNIs (*Calico, Cilium*)
- **Need application-aware network policies (*L7 like HTTP and gRPC*) & mutual encryption:** Proxy-based Service Mesh (*Consul, Istio, Linkerd*)



@sebiwicb



- **Pure service discovery:** Application libraries (*Consul, Netflix Eureka*); CNIs; kube-proxy or Service Mesh
- **Need observability:** CNIs or Service Mesh
- **Need retries, circuit breaking, timeout:** Application libraries (*better customizability*)

... some of the more exotic use cases:

- **Cross-datacenter, cross-cluster network administration:** Cilium or Proxy-based Service Mesh

How about reliability?

- **Proxy-based Service Mesh:** Better support, easier to debug, easier to understand, battle-tested
- **Cilium and others:** Higher performance, good support, but quite new



200  
OK

GitHub to the demo application, slides  
and other materials

